

# Yash Uttamchandani

 yashuttamchandani92@gmail.com ||  +91-9638626569

 github.com/iamyash7 ||  linkedin.com/in/mainhoonyash ||  mainhoonyash.vercel.app

## PROFESSIONAL SUMMARY

Data Scientist & AI/ML Engineer with 2+ years of production experience at scale. Expertise in LLM fine-tuning & prompt engineering, RAG pipelines, NLP, agentic AI, and MLOps on Databricks & Snowflake. Delivered 95%+ reduction in manual effort and 30% uplift in user conversion. IIT Kanpur M.Tech (CS), AIR 185 in GATE CS 2022.

## EDUCATION

| Year      | Degree/Certificate                       | Institute                                          | CPI/%    |
|-----------|------------------------------------------|----------------------------------------------------|----------|
| 2022–2024 | M.Tech / Computer Science & Engineering  | Indian Institute of Technology, Kanpur             | 8.1/10.0 |
| 2019–2022 | B.Tech / Computer Science & Engineering  | Charotar University of Science & Technology, Anand | 8.7/10.0 |
| 2016–2019 | Diploma / Computer Science & Engineering | A.V Parekh Technical Institute, Rajkot             | 8.7/10.0 |

## EXPERIENCE

- Interglobe Aviation (IndiGo)**

Data Scientist

June 2024 – Present  
Gurugram, India

    - Intelligent Feedback Analysis System using LLMs**
      - Engineered production-grade data pipeline (Snowflake → Databricks) with automated daily inference for multi-label sentiment classification (3 sentiment classes + 4-tier issue taxonomy) using few-shot prompt engineering and multi-model evaluation (Claude, LLaMA, Mistral)
      - Operationalized NLP pipeline on Databricks Workflows with structured outputs persisted to Snowflake; built Power BI dashboard reducing manual feedback review effort by 95%+(L1→L2→L3→L4).
    - Hotel Booking Propensity Model using Customer Behavioral Segmentation**
      - Two-stage hierarchical segmentation: Applied Agglomerative Clustering (Ward linkage, Euclidean metric) on 5,000+ confirmed hotel bookers across 20+ behavioral CDP features to identify 16 high-propensity segments; used K-Nearest Neighbors (k=1, distance threshold 1.65) on remaining 40,000+ customers to identify near-bookers behaviorally similar to established bookers vs. distant non-bookers, consolidating 18 clusters into 7 final behavioral cohorts.
      - Targeted audience identification & business impact: Identified 150K+ high-propensity hotel booking candidates segmented by behavioral similarity; enabled marketing team to execute tiered propensity-based campaigns, driving 30% increase in inactive-to-active conversion and 20% uplift in click-to-open rate on targeted hotel promotions
    - LLM-Based Cargo Commodity Code Classification & Revenue Leakage Detection**
      - Built end-to-end LLM pipeline to detect 20% mislabelled commodity codes; implemented batched inference strategy (10 descriptions per LLM call) using LLaMA 4 via Databricks ai\_query(), achieving 95%+ accuracy while reducing API call overhead by 90%
      - Developed rate calculation engine quantifying revenue impact per shipment; automated detection system to Snowflake identified 5% of total cargo revenue exposure from mislabelled commodity codes, enabling proactive correction by ops teams
    - ASK\_PRISM - Conversational Text-to-SQL Intelligence System for Business Self-Serve Analytics**
      - Designed and developed ASK\_PRISM, an enterprise-wide conversational AI system enabling business users to query databases in natural language without requiring SQL expertise or BAU team intervention.
      - Initiated deployment with ASK\_PRISM\_CDP on Databricks Genie Spaces, targeting CDP data for customer and behavioral analytics; successfully deployed to production.
      - Architected an expansion onto Snowflake's Cortex layer — leveraging Cortex Analyst, AI Agents, and Snowflake Intelligence — to cover enterprise-wide data domains including Sales, Air Operations (AIROPS), India Data Hub (IDH), and Catering Operations (COPS).
      - Configured and curated semantic models, table definitions, and domain-specific context for each data domain, enabling accurate natural language to SQL translation across diverse and complex schemas.
      - Collaborated with a cross-functional team to democratize data access across multiple business units, significantly reducing query turnaround time that previously required manual SQL intervention.

## PROJECTS

- Sampling and Reconstruction of Volume Data (CS661) | Guide: Prof. Soumya Dutta**

(Jan 2023 – Apr 2023)

    - Implemented Simple Random Sampling to sample a given percentage of points from the Isabel 3D dataset and performed reconstruction using nearest-neighbor and linear interpolation methods.
    - Attained a signal-to-noise ratio (SNR) of up to 40 across various sampling percentages.

## ACHIEVEMENTS AND CERTIFICATIONS

- Secured **All India Rank 185** in **GATE CS 2022** amongst 77,257 candidates
- Algorithmic Toolbox** — University of California San Diego (Coursera)
- Career Essentials in Generative AI** — Microsoft and LinkedIn

## TECHNICAL SKILLS

- Languages:** Python, SQL, Java, C++, C
- ML / Data Science:** Scikit-learn, NumPy, Pandas, PySpark, Matplotlib, Seaborn, Hugging Face Transformers, MLflow, REST APIs / FastAPI
- Data Engineering:** Snowflake, Databricks, Delta Lake, Delta Live Tables, Apache Spark, Git, AWS
- AI / LLM & Generative AI:** Large Language Models (LLMs), Prompt Engineering, Few-Shot Learning, NLP, Sentiment Analysis, Agentic AI, Multi-Agent Frameworks, Snowflake Cortex, Databricks Genie Spaces